

CS 350 Assignment 1: C and Unix Warmup

Spring 2019

Assigned: January 28, 2019

Due: Always consult Blackboard for due dates.

Overview and Objectives

In Assignment 1, you will write a simple program that exercises several features of the C programming language and of the Unix operating system and programming environment. These features will prove useful for future CS 350 programming assignments this semester, and for subsequent C/Unix systems projects. In particular, you will:

- Write a simple `makefile` that builds an appropriately named executable program and cleans generated files out of a top-level directory. Your `makefile` will support *separate compilation* of code into object files, and will link them together in a separate step into an executable. You may have created similar `makefiles` for CS 240 or other prior classes.
- Write code, or include and call libraries, to traverse and extract command line arguments within a C program.
- Write code to find and use an environment variable.
- Incorporate timing code into a C program.
- Use a random number generator to build input cases.
- Read and write data from and to Unix text files.
- Use `malloc()` to acquire dynamically allocated memory (and `free()` to return it).
- Use Unix pipes (on the command line for now!) to transfer input data from the output of one program to a program that reads from `stdin`.
- Use I/O redirection to send data from a program that writes to `stdout` and `stderr` into two different output files.
- Build flexibility into your program (in this case, for example, by using files or standard input and output for I/O).
- Check all input parameters and return values of system calls; report and handle all errors and invalid arguments appropriately.
- Apply the naming conventions for directories, executables, and submitted archive files that we will use throughout the semester.
- Structure and name your assignment directories and files in the only acceptable CS 350 way, for submission to Blackboard.
- Follow verbose detailed directions!

Specification

You will write two programs that compile into two executables, one for *sorting* integers, and one for *generating* random integers to sort.

Your *generator* program will generate a sequence of random positive integers according to parameters specified on the command line. Your *sorter* program will read a sequence of integers, sort them, and write them out in sorted order. Details of each program appear below.

Sorter

The *sorter* reads a sequence of integers from standard input or from a file and outputs three different things:

- the sorted values, to standard output (`stdout`) by default, or to an output file (if one is specified by the user)
- a report about the counts of a few of the integers in the input, written to standard output by default, or to an output file if one is specified by the user
- the elapsed time of the program, written to standard error (`stderr`).

To sort, you should call the `qsort()` C library function. Read the man pages or online documentation to determine how to prepare the data for `qsort()`, and how to call the function. *Using new library and system call functions is fundamentally important to becoming an effective systems programmer.* Sorted values should appear one per line of output, in text format. Duplicate values in the input should appear in the output.

Your program should count and report the number of occurrences of integers that correspond to the ASCII value of letters in the Linux *userid* of the person running the program. An example report for user *mlewis* appears below.

```
m 109 5
l 108 3
e 101 6
w 119 7
i 105 2
s 115 5
```

The output should appear exactly as above; the *userid* characters (no matter whether they are lowercase letters, capital letters, or digits) appear in the order they appear in the *userid*, one per line. Each character is followed on each line by the ASCII value of the letter and then a count of the number of times that ASCII value appeared in the input. One tab character (`'\t'`) should separate the letter and the ASCII value, and one more should separate the ASCII value and the count. Lines should contain no trailing whitespace, other than a single newline character (`'\n'`). Repeated letters in the *userid* should be repeated in the output. For example:

```
n 110 69
e 101 58
l 108 66
l 108 66
y 121 61
4 52 53
```

Sometimes we will build programs that check your output automatically. If it is not formatted exactly according to specifications, then either we have to take off way too many points, or grading becomes much more difficult. Let's get used to specifying (me) and producing (you) precisely formatted output.

Whether the program uses files or standard input and output depends on command line options, as described in more detail under the program command line options below.

Your program should also report the time elapsed, in seconds (with a decimal part that includes precision to the microsecond) of your program. Timing should be done within your code, not from the shell. Start the timing after parsing command line arguments, and stop the timing after writing the last bit of output (other than time elapsed). Report the elapsed time on `stderr`, not on `stdout`.

The positive integers to be sorted appear as input, one per line, in text format (i.e. not binary). Text format can be read directly by `scanf()` (among other functions), and written using `fprintf()` (among other functions).

You may assume that the input contains valid positive integers only, and ends with a `Control-D`, the Linux end-of-file (EOF) character. Your program should read until EOF or the number of specified integers is reached (see `-n` below), whichever comes first.

Please do *not* prompt the user for values; you should assume that the user knows how to use the program. You may include instructions after printing a "usage string" in response to "`proglsort -u`" if you wish.

Generator

Your generator program writes randomly generated integers to standard output (by default), or to a file (if specified by the user). The integers themselves should be preceded by the number of integers that will follow, should appear one integer per output line, separated only by a single newline character (`'\n'`). This specification is intended to match the input format that the sorter program expects. The number of integers and the range within which they fall are specified on the command line (see below), or with default values (also specified below). Your *generator* program should be compatible with your *sorter* program in the sense that your *sorter* should successfully read and sort integers produced by your *generator*.

Interface, Options, and Parameters

Your programs should support the following interfaces:

```
progl sorter [-u] [-n <num-integers>] [-m <min-int>] [-M <max-int>]  
             [-i <input-file-name>] [-o <output-file-name>] [-c <count-file-name>]
```

```
progl generator [-u] [-n <num-integers>] [-m <min-int>] [-M <max-int>]  
                [-s <seed>] [-o <output-file-name>]
```

Brackets indicate that all the arguments are optional; your program should support a replaceable default for each, as specified and described below. The user may specify options in any order. The meaning of each argument follows:

-u

Print a usage string for your program on `stderr` and then exit. If **-u** appears anywhere on the command line, the program should ignore all other arguments, print the usage string, and exit.

-n <num-integers>

Your *generator* program should generate `<num-integers>` random integers.

Your *sorter* program should read (at most) the first `<num-integers>` values from the input file (or from standard input), sort them, and write them out to the output file (or standard output).

If the input file (or `stdin`) contains more than `<num-integers>` values, your program should simply ignore the rest.

For both programs, the minimum value for `<num-integers>` is 0. The default value for `<num-integers>` is 100. There is no maximum value for `<num-integers>`. If the user specifies a negative value for `<num-integers>`, the program should print an informative error message to `stderr` and exit immediately.

-m <min-int>

Your *generator* program should generate integers no smaller than `<min-int>`.

Your *sorter* program should check for integers in the input that have a value less than `<min-int>`, and halt with an informative error message as soon as the first one appears.

The default and minimum acceptable value for `<min-int>` is 1.

-M <max-int>

Your *generator* program should generate integers no larger than `<max-int>`.

Your *sorter* program should check for integers in the input have a value greater than `<max-int>`, and halt with an informative error message as soon as the first one appears.

The default value for `<max-int>` is 255. The maximum value for `<max-int>` is 1,000,000. `<max-int>` must not be less than `<min-int>`.

-s <seed>

By default, your *generator* program should seed the random number generator with a value derived from reading the system clock. If the user specifies the **-s** option, then your program should instead seed the random number generator with `<seed>`. You may assume that `<seed>` is an unsigned long (and pass it directly to `srand()`). (When writing code that produces pseudo-random values for testing, seeding the random number generator with a time value is good practice. When debugging your code, you often want the same random values to be produced for every run, so that the same behavior occurs each run. This option allows you to select whether the numbers are pseudo-random or deterministic, on the command line.)

Note that the `-s` option is not allowed for the *sorter* program.

`-i <input-file-name>`

By default, your sorter program should read input from `stdin`. If `-i` is specified on the command line, it should instead read input from the file named `<input-file-name>`.

Note that the `-i` option is not supported by the *generator* program.

`-o <output-file-name>`

By default, your programs should write output to `stdout`. If `-o` is specified on the command line, it should instead write output to the file named `<output-file-name>`. Your programs should overwrite the contents of `<output-file-name>` if it exists when your program is invoked.

`-c <count-file-name>`

As described above, your *sorter* should write the ASCII value and a count of the number of instances of each letter in the Unix username of the person running the program. (Do not hardcode your own userid into your program!) By default, your sorter program should write these values and counts to `stdout`. If the user specifies the `-c` option, then your program should write the counts to `<count-file-name>`.

Note that the generator program does not support the `-c` option.

If the user specifies an option that your program does not support (`-z` for example), your program should print the usage string to standard error (`stderr`), and exit. If the user does not specify an argument where one is expected, the program should likewise report the problem, print the usage string to `stderr`, and then exit. You may assume that file names do not begin with a dash, and that the user specifies no negative numbers on the command line. Thus, you may assume that every command line option that begins with a dash (`'-'`) corresponds to an option to your program.

Additional Requirements

In addition to specifying the interface for running your program (above), I will often impose some requirements about *how* you need to implement the assignment. I will try to be clear about “suggestions” vs. “requirements”; when the assignment description does not make the distinction clear, please ask. Requirements will likely appear in the grading rubric.

This assignment’s internal requirements include:

- Any array that you use must be dynamically allocated, using `malloc()` (or some form of `malloc()`).
- Check that all command line input parameters fall within specified bounds. (This is an implicit requirement for all future assignments, unless otherwise stated.)
- Whenever appropriate, check the return value from system calls and library calls, and take appropriate action. When a system call sets `errno` and the error string, you should report that error message to the user. See below for more information about this. (This is an implicit requirement for all future assignments.)

Some Additional Information

You may or may not already know how to parse command line arguments, access environment variables to determine the userid, generate random numbers, time your program from within, allocate memory dynamically, and do the other things that this assignment asks. The text below provides some information about the names of functions and system calls that you should investigate. This is not necessarily an exhaustive list, nor do you necessarily have to use every one of these functions. Sometimes they will point you to other functions that you may prefer, and sometimes you may decide to do things differently.

*Please get used to investigating function and system call usage with **man** and **google**, but please also ask us for help if you need it!*

Parsing command line arguments

You may decide to parse command line arguments “by hand,” or you may decided to learn and use `getopt()`. To get started, you should know that command line arguments are made available to a C programmer through parameters to `main()`.

```
int main(int argc, char **argv) {
```

The parameter `argc` contains a count of the number of command line arguments, and `argv` contains a pointer to an array of pointers to the arguments, each of which is a null terminated character string. Using the man page and Google for `argc`, `argv`, `getopt()`, and `atoi()` or `strtoul()` (to convert character string to an integer or to an `unsigned long`, respectively) should set you on the right path toward parsing command line arguments properly for all of our assignments.

Reading environment variables

Actually, the full interface to `main()` is:

```
int main(int argc, char **argv, char *envp[]) {
```

Note that “`char **p`” and “`char *p[]`” are equivalent. I picked the second form for `envp` just to be different from the `argv` parameter. The userid of the person running your program is held in the environment variable `$USER`. Run the following command to verify that:

```
$ echo $USER
```

Traverse and print the values stored in `envp` until you find the one that includes the user name. To get the value associated with the `$USER` variable, you may want to use `strcmp()` and/or `strtok()`; these functions are cool, but probably overkill in this case. Alternatively, you could simply move a pointer just past the equals sign in the appropriate `envp` array; the rest of that string is the user id you will need. Even better, investigate `getenv()` and avoid this stuff altogether. To exercise your C muscles, dig out `$USER` by hand; for the most elegant code, use `getenv()`. You decide.

Generating random numbers

Investigate `srand()` and `rand()`. Calling `srand()` seeds your random number generator with an initial value. Passing the same integer every time facilitates debugging, because the sequence of random numbers will then be common across multiple runs of the program. After debugging, you should seed the random number generator with a time value read from the system clock. Investigate `gettimeofday()`, which you can also use to time your program.

Timing your program

Investigate the `timeval` struct, `gettimeofday()`, and `difftime()` to learn how to time your program from within. You will need to include `time.h`. Man pages and most (good) information about system and library calls will tell you which header files to include for each function you wish to use. If the compiler complains that it cannot find a function, you may not have included the necessary header file appropriately. If you include your timing code in a separate `.o` file, you will be able to easily include it in all future assignments, whenever necessary. See below under “Separate Compilation...”.

Reading and writing files

Use your favorite way of reading input from text files and `stdin`. I suggest `scanf()` and `fscanf()`. For output, use `fprintf()`.

Allocating memory dynamically

The `malloc()` function takes as an argument the size of the memory being requested. It returns a pointer to that memory. To use `malloc()` properly, you will need to also investigate the `sizeof()` function, and know how to cast a pointer to one type into a pointer to another. Remember to always check the return value from `malloc()` and (just about) every other function you call. If `malloc()`'s return value indicates an error, your program should exit gracefully. Your code is likely to include something like this:

```
int *myints;
myints = (int *) malloc (n * sizeof(int));
/* check the return value here! */
```

For this program, check out the `realloc()` function, which may come in handy because you do not know how many integers may appear in the input stream before you start reading them.

Reporting Error Messages

Some system calls and library functions set a global variable called `errno`. If an error occurs within a system call that sets that variable, you can print a descriptive message by calling the function `perror()`, which takes a pointer to a character string as an argument. It is common, therefore, to include something like the following:

```
if (myints == NULL) {
    perror("Error with myints: ");
    exit(1);
}
```

You may also decide to use the `assert()` function, which can shorten code. Google it.

Code Structure

All source code, header files, and supporting files should be located under a single top-level directory named `ProgNLastname_userid`, where ‘N’ is replaced by the assignment number (in this case 1), “Lastname” is replaced by your last name, and `userid` is your login id. For example, my top-level directory for this assignment would be `Prog1Lewis_mlewis`. This will help us organize 70+ submissions and automate grading.

Place a single `makefile` in your top level directory that builds all object code (that is, “.o files”), libraries, and executables. Unless otherwise specified, please have these generated files reside directly within the top-level directory, *not* within subdirectories. (In other words, even though it may be better practice for larger software systems, please do not build executables into a subdirectory named `Prog1Lewis_mlewis/bin`, for example.)

Your `makefile` should generate all necessary files in response to a simple “make” command, with no command line arguments, executed within the top level directory.

Your `makefile` should remove all executables, `.o` files, and generated libraries, when the user runs “`make clean`”.

Separate Compilation and Executable Naming

To facilitate grading, all of your submissions must have the same name and support the same command line arguments, which I will typically specify for you in assignment descriptions. Your `makefile` should build the executables for this Assignment 1 into files named `progl sorter`, and `progl generator`.

You should use separate compilation to build multiple `.o` files, one per `.c` file. Those `.o` files should then be linked together into an executable in a separate step in your `makefile`.

Your two programs should share some command line parsing code, however you see fit. Each program supports the same arguments for setting maximum and minimum values, so code that does that should not appear in both the sorter and the generator. Instead, it should appear on a `.c` file that both programs use, so you could change a single value (the largest allowable integer, for example) in one place, and have it be reflected in both programs after you have rebuilt them using `make`. (Note, however, that the allowable command line arguments for the sorter and generator are not exactly matched.)

The `-c` option tells `gcc` to generate an object code file, omitting `-c` tells `gcc` to generate an executable, and `gcc` can be invoked with `.o` files, which it then passes to the linker to create an executable. Therefore, your `makefile` should have separate rules for each source file, and one rule per executable.

Playing with Pipes and I/O Redirection

Investigate (again, `man` and Google are your friends!) I/O redirection and pipes, along with the Unix command `cat`. If you have implemented your code properly, you should be able to invoke your generator program, have it write to standard output, and then on the same command line pipe that output to the input of your *sorter* program. Something like this (with command line arguments specified if you want to replace defaults):

```
$ proglgenerator | progl sorter
```

You should also be able to redirect the output of your program to a file... you should be able to do this on the command line, not just within the program using the `-o` option. Try this:

```
$ proglgenerator > myfile.txt
```

Now use `myfile.txt` as the input to another run:

```
$ progl sorter -i myfile.txt
```

You could do the same thing as follows:

```
$ cat myfile.txt | progl sorter
```

Make sure this kind of stuff works; we will use it for testing. One thing that could be a problem is if you have not separated the output that is supposed to be written to `stderr`, from the output to be written to `stdout`. Find out how to redirect `stderr` output to a file.